



About Mercor

What is Mercor?

Mercor turns human intelligence into AI advancement.

Today, the biggest bottleneck to improve AI models isn't compute or engineering talent — it's **high-quality human data**. Models learn their best judgment from people who've done the work. Expert knowledge transfer is a critical engine of AI progress.

We are a trusted partner to AI labs and enterprises to provide expert-written answers, precise feedback, and nuanced evaluations across law, medicine, coding, finance, and literally every other profession and knowledge-domain.

Mercor builds the infrastructure that makes this possible.

We are a human-data platform for frontier AI development. Our platform and operations power the recruiting, onboarding, training, managing, and payments for thousands of experts in hours. Need 150 lawyers to grade legal reasoning tasks? Done by tomorrow. Need 400 software engineers to simulate complex development environments to train AI models? Already ramping.

We've created a **new category of work** that's removed all of the traditional friction of hiring — where professionals in every domain can instantaneously apply and get onboarded onto well-paying projects. We are building closed-loop systems where expert judgment becomes structured signal, signal becomes training data, and outcomes feed back into increasingly precise performance prediction. Over time, this compounds into a proprietary data and tooling layer that directly improves how models reason, plan, and act.

Today, Mercor is a trusted partner to many of the world's leading AI labs. We operate hundreds of concurrent projects spanning evaluation, reinforcement learning, agentic workflows, and post-training.

By creating a platform that sources the highest quality talent in every professional domain, we're now running hundreds of projects that all work towards redefining how AI creates economic value.

Other Mercor facts

- Fastest growing company of all time (\$1M to \$500M revenue run rate in 17 months)
- 30,000+ active experts across hundreds of domains
- \$2M+ paid each day to our talent network
- We partner with all the leading AI labs and 6 out of the Magnificent 7
- Headquartered in San Francisco with a global workforce of more than 300 full-time employees



Engineering & Research

We are currently at the frontier of post-training dataset development, with significant investments in long-horizon evals and RL environments. Our focus is on accurately measuring and improving model performance on complex, multi-step tasks that more closely reflect real-world reasoning in economically valuable domains.

What we build

Platform: Matching, Prediction, and Feedback

At the core of Mercor is a recruiting platform that continuously learns from outcomes.

Our matching and routing engine learns from multiple signals from each applicant: expert background and AI interviews, task-level performance and rubric scores, and past performance. Based on this, we are able to confidently predict future performance, allocate work, and improve data quality over time. We are processing millions of applications and interviews a year.

These systems form closed feedback loops where every completed task sharpens future decisions.

Human data products

On top of this platform, Mercor builds human data products that support the full lifecycle of expert-driven training data.

These systems enable:

- task and workflow generation across hundreds of professional domains
- expert evaluation with rubric-based grading
- reinforcement learning and post-training pipelines
- agentic, long-horizon task execution
- quality control, verification, and performance prediction

Mercor is a pioneer in rubric science and evaluation systems across all modalities: text, image, audio, video, and code. Our platforms run continuously under real customer load and directly influence how frontier models are trained and evaluated.

APEX: Economically meaningful evaluation

APEX is Mercor's AI Productivity Index, designed to evaluate models based on their ability to perform economically valuable tasks across professional services, medicine, software engineering, and consumer activities.

Instead of testing isolated skills, APEX measures whether a model can complete realistic professional workflows using real inputs, sources, and constraints. Tasks are long-context, rubric-graded, and designed by domain experts to reflect how work is actually done.

Engineers and researchers contribute to:

- evaluation frameworks grounded in real-world economics
- scalable grading systems with high consistency
- long-context and multi-step task design
- expansion into new domains, workflows, and environments

APEX is also a foundation for future work in simulated environments where models interact with tools, files, and applications, closing the gap between benchmarks and production reality.

What it's like to build here

- direct exposure to frontier AI labs and post-training workflows
- rare insight into how models actually improve in production
- unprecedented capital, growth, and scope of ownership
- hard problems at the edge of current AI capability

Engineering here means shaping how human intelligence becomes machine-learnable at scale.

Equity & Benefits

Pricing

- **Preferred price/share:**
 - We calculate the value of your grant by multiplying the number of shares you've received by the preferred price per share at our most recent fundraising round. Please note that there is no guarantee that you will ultimately be able to sell your shares at the preferred price per share.
- **Stock options strike price:**
 - This will be the company's most recent 409A valuation, as determined by a third-party valuation firm.
 - This is the price you pay to exercise your options and receive your shares.
 - The date of options exercise starts the clock for long-term capital gains tax treatment.

Company benchmarks for market capitalization

When you're thinking about the potential upside of your Mercor equity, it helps to see how capital-markets value similar businesses.

The four companies below were chosen because they either (a) own large hiring or professional-network platforms, or (b) sell software that sits at the center of recruiting decisions. As of mid-2025:

- LinkedIn is widely modeled by sell-side analysts at \$150Bn were it spun out of Microsoft.



- Recruit Holdings (parent of Indeed & Glassdoor) trades publicly at \$78Bn despite sub-12% net margins.
- Scale AI last raised at a \$29Bn valuation.
- Rippling last raised at a \$16.8Bn valuation.

We're already operating at >30% gross margins, more than triple Recruit's and well above most private-market peers. The takeaway is that capital-efficient, high-margin businesses in our space easily command multi-billion-dollar outcomes.

Please note that the preceding information is being used for demonstration purposes only and we cannot guarantee similar results.

Benefits

We offer the following benefits:

- Anthem Gold PPO E healthcare plan with Mercor covering 100% of the premium for the employee
 - You have the option to upgrade to the Anthem Platinum PPO A plan and pay the difference
- Kaiser HMO healthcare plan with Mercor covering 100% of the premium for both employees and dependents
- Vision & Dental coverage through Guardian
- Flexible benefits: FSA, HSA (with enrollment in HDHP), & Commuter Benefits
- \$1,500 per month meal stipend while working in the SF office
- \$10,000 per year proximity bonus for living within 0.5m of the office
- Up to \$300/month for a gym membership at Equinox Sports Club San Francisco
- 3 months of fully-paid parental leave if you are the birthing parent and 2 months of fully-paid parental leave if you are a non-birthing parent
- Unlimited PTO, subject to manager approval

FAQ

What problem does Mercor solve?

Mercor exists to turn human intelligence into AI advancement.

As AI models improve, the primary bottleneck is no longer compute or architecture, but high-quality human data. Modern models need expert judgment to learn how to reason, evaluate, and act in real professional contexts.

Mercor builds the infrastructure that turns human expertise into structured, model-ready data at scale. We enable AI labs to fine-tune, evaluate, and post-train models using reliable, expert-generated signals across many domains and workflows.

What are work expectations?

We are an ambitious team, and our commitment to our customers sets a fast pace. Success in this hyper-growth environment requires a level of dedication that goes beyond



a traditional 9-to-5. This is a significant opportunity for individuals who thrive on this kind of challenge and are ready to accelerate their career.

We work in-person Monday through Friday and all team members start at 9 am. Similar to all start-ups, there is typically weekend work that people complete remotely, though many team members choose to come in on weekends.

What is human data?

Human data is the carefully judged feedback that real people (doctors, lawyers, engineers, teachers, and other specialists) give to an AI system so it can learn what “right” looks like. Instead of letting a model guess from random internet text, we show it trusted answers and corrections from subject-matter experts, much the way an apprentice learns by watching a master and receiving pointers. Over time, those expert signals shape the model’s judgment, helping it tackle harder, more professional tasks with fewer mistakes.

What does human data look like in the long-term?

Human data is already a multi-billion-dollar market, and Mercor is on track to become its dominant supplier. Each time model capability advances, it creates new problems that a model cannot evaluate by itself. As soon as today’s benchmarks fail, researchers publish harder ones, covering long-context reasoning, multi-step workflows, and application use so demand for expert evaluators grows with progress.

The shift to reinforcement learning (RL) based approaches necessitates high volumes of human-generated evaluation data (evals). Collect enough of those rewards and the model is able to climb the competency ladder itself. That is why human-generated feedback will be the lever that unlocks the next decade of progress.

Is Mercor at the frontier of AI progress?

Mercor partners with, and learns from, every top research lab and application builder in the field. Mercor sits at the intersection of their experiments: our team sees the freshest model capabilities, post-training techniques, and deployment pain-points weeks or months before the broader market. That bird’s-eye perspective guides the products we build and lets employees work on what could become one of the most valuable data and tooling layers in the AI stack.

Is Mercor a “services” company?

Mercor is best understood as software-defined human data infrastructure.

Historically, services companies had structurally lower margins than software companies because delivery scaled linearly with people. In the AI era, software companies increasingly carry high ongoing costs from compute and inference, while modern, AI-native platforms can automate large portions of what used to be manual services work.

Mercor sits squarely in this new category. While our platform coordinates large numbers of human experts, the value is created by software: automation, evaluation systems,



quality controls, routing, and performance prediction. These systems allow human work to scale non-linearly, with far fewer costs below gross margin than traditional services models.

Under the hood, Mercor operates a deep software stack that would, in many cases, stand alone as major products:

- Payroll rails
- AI interviewer & assessments
- Listings & ATS
- Auto-sourcing search
- Document ingestion/parsing, mock interviews, upskilling

Where is the office located?

Our office is located on the 33rd floor of 181 Fremont in San Francisco, CA 94105 and we plan long-term to be based in the SOMA District in San Francisco.

Why is now the right moment to join?

The product is live, profitable, and scaling fast, but the org chart is still flat enough that individual decisions visibly move revenue. Joining now means influencing culture and architecture before we add the next few hundred people.

Visiting Mercor

This guide covers what to expect when you reach the onsite stage of your interview with Mercor, including how to prepare, what your day will look like, and travel details if applicable. We've outlined everything you need to ensure a smooth and comfortable experience. If you have any questions, don't hesitate to reach out.

Preparing for Your Visit

Timing and format

Interviews are typically scheduled between 12:00–6:00 PM, Tuesday–Friday, and are usually booked about two weeks in advance. Earlier scheduling can often be accommodated if needed. Please provide as much availability as possible—longer blocks are preferred.

If you're based in the Bay Area, interviews will take place in person at our office. Otherwise, interviews will be conducted virtually.

Tools and preparation

AI tools such as Cursor or an LLM may be used unless otherwise specified. If applicable, please set up access in advance. Any required API-related expenses may be reimbursed.

NDA

An NDA will be sent prior to your visit. Please complete it before arrival.



Arriving at Mercor HQ

Location

Mercor HQ
181 Fremont Street, San Francisco, CA 94105
33rd floor

Check-in

Plan to arrive 15 minutes before your first interview. A government-issued photo ID is required for building access. You will be on the guest list—please check in with security in the lobby before proceeding upstairs.

Your Time Onsite

Schedule

Expect to be onsite for approximately 4–6 hours, including breaks. You will receive a calendar invite with your full schedule. All times are listed in Pacific Time unless noted otherwise.

Work setup

We'll provide a company laptop for all interview sessions. You're welcome to bring your own device for reference, but all work should be completed on the provided laptop unless noted otherwise.

Environment

The dress code is casual—wear what's most comfortable.

During your visit, you may be placed in a shared workspace with other candidates. We're growing quickly and space is limited, but we'll do our best to ensure you're comfortable throughout the day.

If you need any accommodations, please let us know in advance.

Meals and refreshments

Lunch during your onsite is reimbursable—please retain your receipt.

Snacks and drinks are available throughout the day at the office, including fresh fruit, protein bars, granola bars, nuts, trail mix, chips, crackers, chocolate, and more.

We also keep a variety of beverages stocked, including still and sparkling water, coconut water, cold brew, coffee, tea, Poppi, kombucha, juice, soda, and other drinks.

Travel & Expenses

Travel

If you are traveling, please book your own flights and accommodations based on what is most convenient and comfortable for you. We do not arrange travel on your behalf. Up to \$1,500 in domestic travel expenses will be reimbursed.



If traveling from outside the Bay Area, we recommend arriving the evening before your onsite.

Reimbursement process

Before your visit, submit your estimated travel costs using the reimbursement form. After your visit, you'll receive a Brex invitation to submit expenses.

Please submit all expenses within 7 days and include itemized receipts and full itineraries. Receipts must be fully processed and include the date, vendor, description, and total amount.

Reimbursements may be denied if documentation is incomplete or does not meet policy requirements.